

This notebook is an exercise in the [Pandas](#) course. You can reference the tutorial at [this link](#).

Introduction

In this set of exercises we will work with the [Wine Reviews dataset](#).

Run the following cell to load your data and some utility functions (including code to check your answers).

```
In [ ]: import pandas as pd

reviews = pd.read_csv("../input/wine-reviews/winemag-data-130k-v2.csv",
index_col=0)
pd.set_option("display.max_rows", 5)

from learntools.core import binder; binder.bind(globals())
from learntools.pandas.indexing_selecting_and_assigning import *
print("Setup complete.")
```

Look at an overview of your data by running the following line.

```
In [ ]: reviews.head()
```

Exercises

1.

Select the `description` column from `reviews` and assign the result to the variable `desc` .

```
In [ ]: # Your code here
desc = reviews.description
# Check your answer
q1.check()
```

Follow-up question: what type of object is `desc` ? If you're not sure, you can check by calling Python's `type` function: `type(desc)` .

```
In [ ]: #q1.hint()
#q1.solution()
```

2.

Select the first value from the `description` column of `reviews` , assigning it to variable `first_description` .

```
In [ ]: # Your code here
first_description = desc[0]
# Check your answer
q2.check()
first_description
```

```
In [ ]: #q2.hint()
#q2.solution()
```

3.

Select the first row of data (the first record) from `reviews` , assigning it to the variable `first_row` .

```
In [ ]: first_row = reviews.iloc[0]
```

```
# Check your answer  
q3.check()  
first_row
```

```
In [ ]: #q3.hint()  
#q3.solution()
```

4.

Select the first 10 values from the `description` column in `reviews`, assigning the result to variable `first_descriptions`.

Hint: format your output as a pandas Series.

```
In [ ]: first_descriptions = reviews.description[0:10]
```

```
# Check your answer  
q4.check()  
first_descriptions
```

```
In [ ]: #q4.hint()  
#q4.solution()
```

5.

Select the records with index labels `1`, `2`, `3`, `5`, and `8`, assigning the result to the variable `sample_reviews`.

In other words, generate the following DataFrame:

	country	description	designation	points	price	province	region_1	region_2	taster_name	taster_twitter_handle
1	Portugal	This is ripe and fruity, a wine that is smooth...	Avidagos	87	15.0	Douro	NaN	NaN	Roger Voss	@vossroger
2	US	Tart and snappy, the flavors of lime flesh and...	NaN	87	14.0	Oregon	Willamette Valley	Willamette Valley	Paul Gregutt	@paulgwine
3	US	Pineapple rind, lemon pith and orange blossom ...	Reserve Late Harvest	87	13.0	Michigan	Lake Michigan Shore	NaN	Alexander Peartree	NaN
5	Spain	Blackberry and raspberry aromas show a typical...	Ars In Vitro	87	15.0	Northern Spain	Navarra	NaN	Michael Schachner	@wineschach
8	Germany	Savory dried thyme notes accent sunnier flavor...	Shine	87	12.0	Rheinhessen	NaN	NaN	Anna Lee C. Iijima	NaN

```
In [ ]: indices = [1, 2, 3, 5, 8]
sample_reviews = reviews.loc[indices]

# Check your answer
q5.check()
sample_reviews
```

```
In [ ]: #q5.hint()
#q5.solution()
```

6.

Create a variable `df` containing the `country`, `province`, `region_1`, and `region_2` columns of the records with the index labels `0`, `1`, `10`, and `100`. In other words, generate the following DataFrame:

	country	province	region_1	region_2
0	Italy	Sicily & Sardinia	Etna	NaN
1	Portugal	Douro	NaN	NaN
10	US	California	Napa Valley	Napa
100	US	New York	Finger Lakes	Finger Lakes

```
In [ ]: indices=[0,1,10,100]
        cols=["country","province","region_1","region_2"]
        df = reviews.loc[indices,cols]
        # Check your answer
        q6.check()
        df
```

```
In [ ]: #q6.hint()
        #q6.solution()
```

7.

Create a variable `df` containing the `country` and `variety` columns of the first 100 records.

Hint: you may use `loc` or `iloc`. When working on the answer this question and the several of the ones that follow, keep the following "gotcha" described in the tutorial:

`iloc` uses the Python stdlib indexing scheme, where the first element of the range is included and the last one excluded. `loc`, meanwhile, indexes inclusively.

This is particularly confusing when the DataFrame index is a simple numerical list, e.g. `0, ..., 1000`. In this case `df.iloc[0:1000]` will return 1000 entries, while `df.loc[0:1000]` return 1001 of them! To get 1000 elements using `loc`, you will need to go one lower and ask for `df.loc[0:999]`.

```
In [ ]: cols=["country","variety"]
df=reviews.loc[0:99,cols]
# Check your answer
q7.check()
df
```

```
In [ ]: #q7.hint()
#q7.solution()
```

8.

Create a DataFrame `italian_wines` containing reviews of wines made in `Italy`. Hint: `reviews.country` equals what?

```
In [ ]: italian_wines = reviews.loc[(reviews.country=="Italy")]

# Check your answer
q8.check()
```

```
In [ ]: #q8.hint()
#q8.solution()
```

9.

Create a DataFrame `top_oceania_wines` containing all reviews with at least 95 points (out of 100) for wines from Australia or New Zealand.

```
In [ ]: top_oceania_wines = reviews.loc[(reviews.country.isin(["Australia","New
Zealand"]))& (reviews.points>=95)] # Check your answer
q9.check()
top_oceania_wines

# Check your answer
```

```
q9.check()  
top_oceania_wines
```

```
In [ ]: #q9.hint()  
        #q9.solution()
```

Keep going

Move on to learn about [summary functions and maps](#).

Have questions or comments? Visit the [Learn Discussion forum](#) to chat with other Learners.